


Mobile & Ubiquitous Computing, 2020 Spring

FedEx: Scalable Framework for Mobile User-Adaptive Deep Learning Training

{Ahnjae Shin, Kyunggeun Lee, Heeseung Yun, Taebum Kim} @ Team 1

Recap

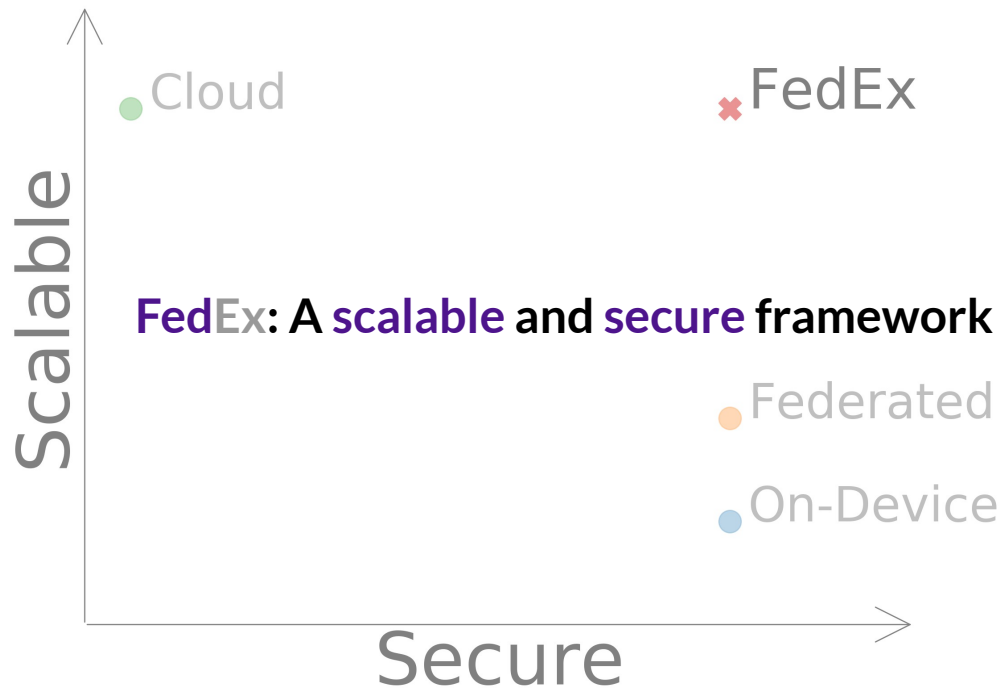


Scalable

Able to scale to large models

Secure

Able to train without raw data






We ship fresh models daily

Federated Learning

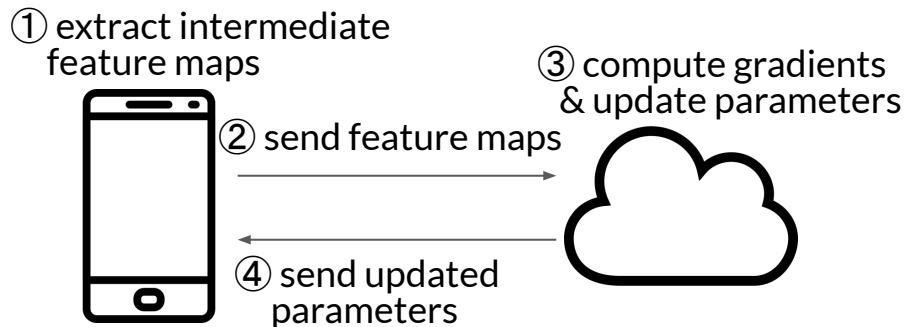
FedAvg

- Compute in mobile devices, update aggregated gradients in clouds
- Computation: mobile device
- Communication rounds: $O(NP)$



FedEx


- Extract intermediate feature maps in mobile devices, train in clouds
- Computation: cloud
- Communication rounds: constant



Federated Learning

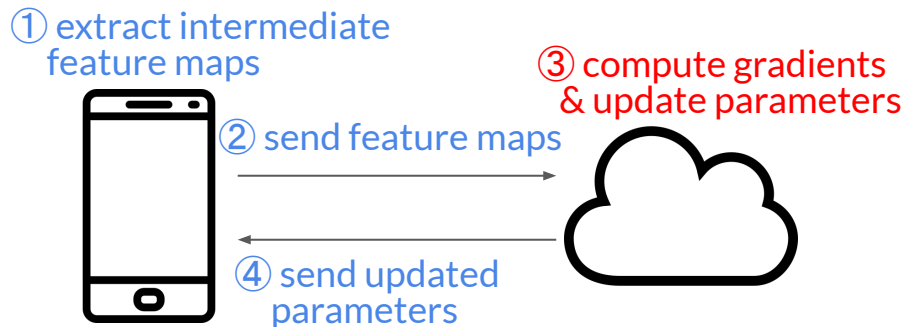
FedAvg

- Compute in mobile devices, update aggregated gradients in clouds
- Computation: mobile device
- Communication rounds: $O(NP)$



FedEx

- Extract intermediate feature maps in mobile devices, train in clouds
- Computation: cloud
- Communication rounds: constant





Refinement

Refinement

- Problem statement
 - Existing federated learning (ex. FedAvg) is slow
 - repeated computation in mobile devices
 - heavy communication between mobile devices and clouds.
- Project main idea
 - Extract intermediate feature maps in mobile devices, and train in cloud servers.
 - Only constant computation in mobile devices.
 - No raw data is sent to the server.
- Project scope
 - FedEx outperforms FedAvg for a federated transfer learning task w.r.t. training throughput.
- **Target application (changed)**

Target application

- Handwriting → **Face recognition**
 - Motivation
 - Necessity of additional training
 - No sequential modeling, moderate complexity
 - More diverse distribution across a variety of users



Rita El
Khoury


Scott
Scrivens

David
Ruddock

Target application

- Handwriting → **Face recognition**

- After literature review, we decided to focus on a more intriguing problem
- Some open questions in face recognition domain
 - Technology of the moment: [masked face recognition](#)
 - Unsupervised face clustering
(ex. How many characters are there in the smartphone gallery?)



Wang et al. (2020). Masked Face Recognition Dataset and Application. *arXiv preprint arXiv:2003.09093*.

Tapaswi et al. (2019). Video Face Clustering with Unknown Number of Clusters. In ICCV.

Target application


- Handwriting → **Face recognition**

- After literature review, we decided to focus on a more intriguing problem
- Some open questions in face recognition domain
 - Technology of the moment: [masked face recognition](#)
 - Unsupervised face clustering
(ex. How many characters are there in the smartphone gallery?)
- Objective
 - For given set of facial images from a few characters, train a model that can discern “Who is who” by leveraging user-adaptive training framework of FedEx

Architecture

System Architecture

Mobile




Cloud




System Architecture

Mobile




Cloud




System Architecture

Mobile



Cloud



Project Status

- Features

- Basic application with TFLite models
- Deploy TFLite models that output intermediate feature maps
- Partial training using intermediate feature maps
- Client-server communication (feature maps & parameters)
- Face recognition application
- Personalized data collection in app

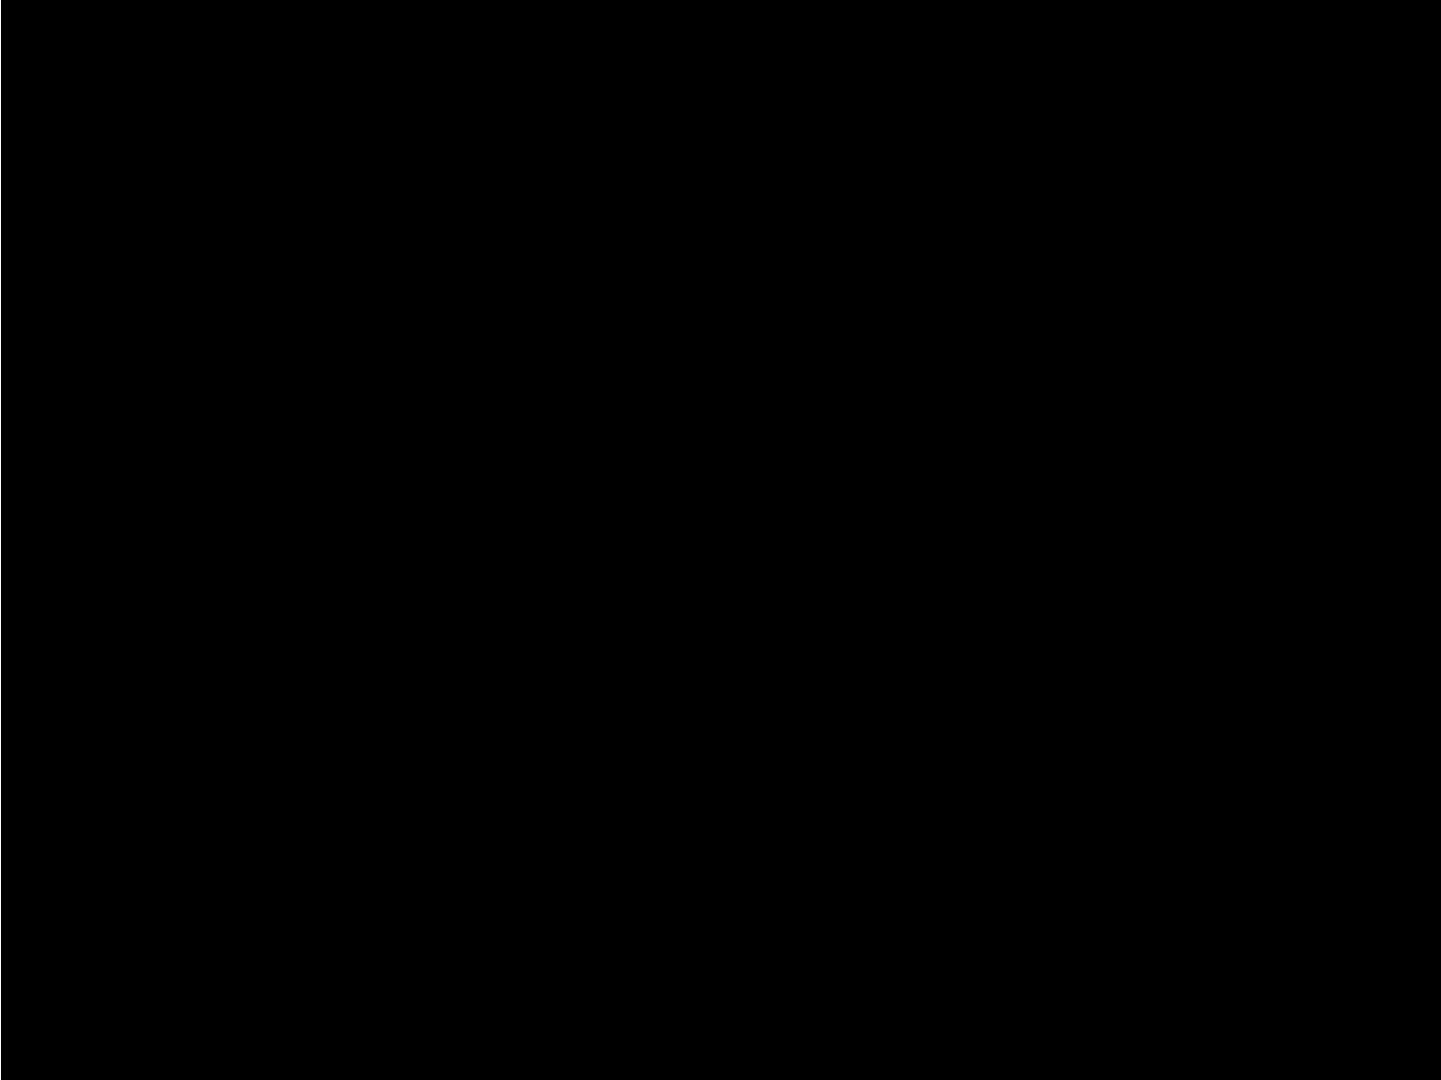
- Evaluation

- Deploy light model
- Deploy heavier models
- Implement FedAvg
- Measure training throughput
- Measure device energy consumption

Demo

Demo

1. Run app that inference DL models
2. The app will save **intermediate outputs as file**
3. Use adb (Android debug bridge) to copy files to cloud
4. Use intermediate outputs to **train model**



Remaining Features

1. Run an app with DL inference models
2. The app will ~~save intermediate outputs as file~~ **send files to cloud**
3. ~~Use adb (Android debug bridge) to copy files to cloud~~
4. Use intermediate outputs to **train model**
5. **Pull new model from cloud**

Challenges & Solutions

Remaining Challenges




- Personalized data collection
 - Should prove that additional training improves the accuracy.
 - However, a significant amount of labeled data is required.
- Implementation of baseline system (FedAvg)
 - The current DL frameworks (TensorFlow, PyTorch, etc.) do not support Python APIs for federated learning.
- Large scale experiment
 - As the number of involved mobile devices grows, FedEx is expected to outperform FedAvg by a larger gap.
 - In reality, however, it is difficult to conduct a very large-scale experiment with a great number of mobile devices.

Possible Solutions

- Personalized data collection
 - Tons of massive annotated datasets available: CelebA, VGGFace2, TrillionPairs, etc.
 - Split the existing dataset for:
 - Primary training
 - Secondary training (sets of multiple identities)
 - It is also possible to leverage YouTube videos by applying face tracking
- Implementation of baseline system (FedAvg)
 - On-device training using Tensorflow Java APIs
 - Simulate mobile devices on desktops (TensorFlow Federated)
- Large scale experiment
 - Simulate mobile devices on desktops (mocking behavior)

Schedule

Overall Plan

#Iter.	Objective	Duration	Misc.
	Ideation & Design Brainstorming, Literature Review, Proposal Feedback	4/1 - 4/14	4/6: Proposal
	System & Architecture Cloud Setup, System Design, Client Mockup	4/15 - 4/28	
	Model & Algorithm Data Preparation, Model Implementation, Integration	4/29 - 5/12	5/11: Project Demo
4	Optimization & Experiment Ablation Studies, Alpha Deployment, Load Testing	5/13 - 5/26	
5	Deployment & Deliverable	5/27 - 6/8	6/8: Final Presentation

Task Assignment

Task	Ahnjae Shin	Kyunggeun Lee	Heeseung Yun	Taebum Kim
Simple client app implementation	O		O	
Server basic implementation		O		O
Light-weight TFLite model implementation & deployment		O	O	O
Face recognition model development	O		O	O
Face recognition app implementation		O	O	O
Baseline solution (FedAvg) implementation & experiment	O	O		
Integrate server and client	O	O	O	O

Deliverable

Deliverable

- Final Deliverable
 - Technical report regarding evaluations
 - Demo application
- Success Criteria
 - Our method succeeds to **train larger model faster** than baseline method
 - Our method succeeds to **use less energy** than baseline method
- Potential Threats
 - Absence/lack of TFLite Android API
 - *Procrastination*

Lesson Learned

- Android application development
 - Gradle build configuration is hard — Android Studio + Gradle seems to have an OS-dependent issue (in Ubuntu).
 - Android application programming/debugging process
- TFLite deployment
 - TFLite converter has become much more powerful and convenient in TF2.x.
 - Yet, subtle issues seem to be still present in TFLite converter (accuracy problem).
- On-device training
 - The field of on-device (or federated) training is still very unexplored.
 - Difficulties in finding reference materials
 - Our project is expected to run into many unexpected issues implementing on-device training.



We ship fresh models daily

Thank you

Backup Slides

Remaining Todos

1. Features

- a. Device-cloud communication: save/load models and features automatically
- b. Label collection: collect personalized labels in our app

2. Evaluation

- a. Implement & evaluate on relatively heavier models
(currently only implemented a light-weight sample model (MobileNet) & task)
- b. Implement or simulate baseline solution (FedAvg)